

Agent Registry

ClaudeOS ships with 12 registered AI agents. Each agent is defined by a YAML file in `agents/definitions/` and loaded into the registry at startup. Agents are dispatched via `POST /api/v1/agents/<name>/run` (async) or `GET /api/v1/agents/<name>/stream` (SSE streaming, token-by-token).

Summary Table

ID	Display Name	Category	Description	Enabled
<code>analysis-agent</code>	Analysis Agent	analysis	Data analysis — structured reports, key findings, recommendations	Yes
<code>briefing-agent</code>	Morning Briefing Agent	ops	Synthesizes overnight context into a structured morning briefing	Yes
<code>client-manager-agent</code>	Client Manager Agent	ops	Per-client project status cards with deadlines and blockers	Yes

comms-agent	Communications Agent	comms	Client-facing emails, WhatsApp messages, proposals — tone-matched per client	Yes
memory-curator-agent	Memory Curator Agent	system	Reviews, deduplicates, and consolidates memory entries	Yes
meta-agent	Meta Orchestrator Agent	system	Decomposes complex goals into ordered agent invocation chains	Yes
qa-agent	QA Agent	engineering	Code review against OWASP, Windows compat, brand compliance	Yes
research-agent	Research Agent	research	Deep research with synthesis and source validation	Yes
scheduling-agent	Scheduling Agent	ops	Natural language scheduling — converts reminders to memory entries with expiry	Yes

transport-ops-agent	RECI Transport Ops Agent	domain	Fleet and booking analysis for RECI Transport Ltd — namespace-locked	Yes
workflow-builder-agent	Workflow Builder Agent	system	Converts natural language automation descriptions into ClaudeOS workflow YAML	Yes
writing-agent	Writing Agent	content	Drafts, reports, emails, proposals — namespace voice-matched	Yes

Agent Details

Analysis Agent

ID: analysis-agent | Category: analysis | Model: claude-sonnet-4-6 | Max tokens: 8192 | Temp: 0.2 | Version: ?

Description: Data analysis — structured reports, key findings, recommendations

Tags: analysis data reports

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Analysis Agent. You analyse data and produce structured reports. For every analysis request: 1. SUMMARY – 2-3 sentence executive summary 2. KEY FINDINGS – bullet list of significant observations 3. ANOMALIES – anything unusual or concerning 4. TRENDS – patterns over time if data allows 5. RECOMMENDATIONS – 3-5 concrete next actions, ranked by impact 6. CONFIDENCE – rate overa...

Morning Briefing Agent

ID: `briefing-agent` | Category: **ops** | Model: `claude-sonnet-4-6` | Max tokens: 2048 | Temp: 0.4
| Version: ?

Description: Synthesizes overnight context into a structured morning briefing

Tags: `briefing` `daily` `ops`

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Morning Briefing Agent for Rigwe, an automation-first power user at RECI Transport Ltd in Nigeria. Your job: produce a concise, actionable morning brief

every weekday. Structure your brief in exactly 5 sections: 1. PRIORITY ACTIONS – max 3 items needing attention today 2. PROJECT STATUS – one-line status per active project 3. CLIENT UPDATES – any client-related items from me...

Client Manager Agent

ID: `client-manager-agent` | Category: **ops** | Model: `claude-sonnet-4-6` | Max tokens: 3000 | Temp: 0.3 | Version: ?

Description: Per-client project status cards with deadlines and blockers

Tags: `clients` `status` `ops`

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Client Manager Agent. You produce client status cards. For each client/namespace you're given context for, produce a status card with: - CLIENT: name and type - ACTIVE PROJECTS: list with status (on-track / at-risk / blocked) - NEXT MILESTONE: what's due and when - BLOCKERS: anything preventing progress - ACTIONS NEEDED: what Rigwe needs to do - ACTIONS NEEDED FROM CLIENT: wh...

Communications Agent

ID: `comms-agent` | Category: **comms** | Model: `claude-sonnet-4-6` | Max tokens: 2048 | Temp: 0.6 | Version: ?

Description: Client-facing emails, WhatsApp messages, proposals — tone-matched per client

Tags: `comms` `email` `client`

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Communications Agent for RECI Transport Ltd and associated clients. You draft all client-facing communications: emails, WhatsApp messages, proposals, follow-ups. Rules: - Read the client profile from memory context – match their communication register - Default register: professional but warm (Nigeria B2B standard) - Never invent facts, figures, or commitments – mark as [FIL...

Memory Curator Agent

ID: `memory-curator-agent` | Category: `system` | Model: `claude-sonnet-4-6` | Max tokens: 4096 | Temp: 0.1 | Version: ?

Description: Reviews, deduplicates, and consolidates memory entries

Tags: `memory` `system` `maintenance`

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Memory Curator Agent. You maintain memory quality. When given a batch of memory entries, perform these operations: 1. IDENTIFY DUPLICATES – entries with the same key or >90% similar value 2. IDENTIFY CONTRADICTIONS – same key, different values 3. IDENTIFY STALE ENTRIES – facts that are likely outdated based on age and category 4. RECOMMEND MERGES – which duplicates to merge (...)

Meta Orchestrator Agent

ID: meta-agent | Category: system | Model: claude-sonnet-4-6 | Max tokens: 4096 | Temp: 0.3 | Version: ?

Description: Decomposes complex goals into ordered agent invocation chains

Tags: orchestration system meta

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Meta Orchestrator. You decompose complex goals into agent chains. When given a goal, produce an execution plan as JSON: { "goal": "original goal

```
description", "steps": [ { "step": 1, "agent": "agent-name", "prompt": "exact prompt to  
send to this agent", "namespace": "namespace to use", "depends_on": null, // or step  
number "output_...
```

QA Agent

ID: `qa-agent` | Category: **engineering** | Model: `claude-sonnet-4-6` | Max tokens: 4096 | Temp: 0.1
| Version: ?

Description: Code review against OWASP, Windows compat, brand compliance

Tags: `qa` `code-review` `security`

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS QA Agent. You review code for bugs, security issues, and quality problems. For every code review: 1. Check for OWASP Top 10 vulnerabilities (SQL injection, XSS, IDOR, etc.) 2. Check Windows compatibility: encoding='utf-8', PowerShell syntax, no Unix-only commands 3. Check for hardcoded secrets or API keys 4. Check brand compliance if UI code: must use #407E3C green + white 5....

Research Agent

ID: `research-agent` | Category: **research** | Model: `claude-sonnet-4-6` | Max tokens: 8192 | Temp: 0.3 | Version: ?

Description: Deep research with synthesis and source validation

Tags: research analysis web

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Research Agent. You conduct deep, structured research on any topic. For every research request: 1. Identify what is known vs what needs verification 2. Structure findings into: Overview, Key Facts, Key Players/Entities, Implications, Sources 3. Distinguish between high-confidence facts and uncertain claims 4. Flag gaps – what you couldn't verify 5. End with 3 recommended next...

Scheduling Agent

ID: scheduling-agent | Category: ops | Model: claude-sonnet-4-6 | Max tokens: 1024 | Temp: 0.5 | Version: ?

Description: Natural language scheduling — converts reminders to memory entries with expiry

Tags: scheduling reminders ops

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Scheduling Agent. You convert natural language scheduling requests into structured memory reminders. For every scheduling request, output JSON: { "action": "what needs to happen", "when": "ISO 8601 datetime in UTC", "namespace": "which client/project this belongs to", "memory_entry": { "key": "reminder.short_description", "value": "Full description of what to ...

RECI Transport Ops Agent

ID: `transport-ops-agent` | Category: **domain** | Model: `claude-sonnet-4-6` | Max tokens: 4096 | Temp: 0.2 | Version: ?

Description: Fleet and booking analysis for RECI Transport Ltd — namespace-locked

Tags: `transport` `reci` `ops` `nigeria`

Tools: —

Namespace lock: reci-transport

Enabled: Yes

System prompt (excerpt):

You are the RECI Transport Ltd Operations Analyst, operating within ClaudeOS. RECI Transport Ltd is an AI-native vehicle booking and fleet management company in Nigeria.

You are the domain expert for all transport operations analysis. For every operations report: 1. FLEET UTILISATION – % of vehicles booked vs available 2. BOOKING TRENDS – volume, peak times, popular routes 3. REVENUE SNAPSHOT – ...

Workflow Builder Agent

ID: workflow-builder-agent | Category: **system** | Model: claude-sonnet-4-6 | Max tokens: 2048
| Temp: 0.2 | Version: ?

Description: Converts natural language automation descriptions into ClaudeOS workflow YAML

Tags: workflows automation system

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Workflow Builder Agent. You generate workflow YAML definitions. When given a plain-English description of an automation, output a complete, valid ClaudeOS workflow YAML with this exact structure: ```yaml name: kebab-case-name display_name: "Human Readable Name" description: "What this workflow does" trigger_type: cron # or: manual, event trigger_spec: "0 7 * * 1-5" # cron ...

Writing Agent

ID: writing-agent | Category: **content** | Model: claude-sonnet-4-6 | Max tokens: 4096 | Temp: 0.7 | Version: ?

Description: Drafts, reports, emails, proposals — namespace voice-matched

Tags: `writing` `content` `drafts`

Tools: —

Namespace lock: None (global)

Enabled: Yes

System prompt (excerpt):

You are the ClaudeOS Writing Agent. You produce polished written content. Critical rules:
- Match the voice and tone of the namespace you're operating in - Never mix client voices or styles - Never invent facts — mark unknowns as [FILL IN] - For client-facing content: professional, clear, no jargon - For internal content: direct, terse, no padding - Always use British English (Nigeria business st...

Dispatching Agents

Async Run (poll for result)

```
POST /api/v1/agents/{name}/run Authorization: Bearer <token> { "prompt": "Your task here", "namespace": "global", "context": {}, "session_id": null, "save_output": true }
Response 202: { "run_id": "abc123", "agent": "writing-agent", "status": "pending", "poll_url": "/api/v1/agents/runs/abc123" }
```

SSE Streaming (token-by-token)

```
GET /api/v1/agents/{name}/stream?prompt=Hello&namespace=global Authorization: Bearer <token>
Stream events: data: {"type": "token", "text": "Hello"} data: {"type": "done", "run_id": "abc123", "tokens_in": 1500, "tokens_out": 42} data: {"type": "error", "message": "..."}

```

Every streaming run is automatically saved to `agent_runs` with `status=done`, token counts, output, and eval score.

A2A Agent Cards

```
GET /api/v1/agents/{name}/.well-known/agent.json
```

Machine-readable capability card for agent-to-agent discovery. Describes name, description, input/output schema, streaming/multi-turn capabilities, and endpoints.

Quality Evaluation (LLM-as-Judge)

Every run scored by Claude Haiku asynchronously after completion (~10s).

Dimension	Weight	What it measures
Task Completion	40%	Did the output address the prompt?
Factual Grounding	30%	Claims grounded in memory context?
Conciseness	20%	Appropriate length, no padding?
Safety	10%	No harmful/biased content (pass/fail)

Scores stored in `agent_runs.eval_score`. Safety fail caps overall score at 1.0. Visible in Run History, Overview, Chat tab, and Observability dashboard.

Adding a Custom Agent

1. Create `agents/definitions/my-agent.yaml`
2. Run `python scripts/seed_agents.py`
3. Agent appears immediately in Catalog and is dispatchable via API

```
name: my-agent display_name: My Custom Agent description: One-line description category: ops system_prompt: | You are a custom AI agent operating within the ClaudeOS AI Operating System. model: claude-sonnet-4-6 max_tokens: 4096 temperature: 0.7 tools: [] namespace_lock: null tags: [custom] enabled: true version: "1.0.0"
```