

ClaudeOS

AI OPERATING SYSTEM

Client Guide & Reference Manual

Version 1.0.0 - May 2026

This document covers everything you need to know about ClaudeOS -- from opening the dashboard for the first time, to running automated workflows and managing your client workspace.

This guide includes:

- Step-by-step instructions for using the dashboard
- Full reference for all 12 AI agents and 7 automated workflows
- Client-specific sections: RECI Transport, Ivycandy Hair, Faiyke AI, Personal
- Memory system, outputs, and cloud sync guides
- Technical reference for administrators and troubleshooting

Table of Contents

1. What is ClaudeOS?	1
2. Design Principles	2
3. Getting Started	4
4. Dashboard Walkthrough	6
5. AI Agents Reference	10
6. Workflows Reference	14
7. Client Workspace Sections	17
8. Memory System Guide	20
9. Outputs Guide	22
10. Technical Reference (Administrators)	24
11. Troubleshooting	26
12. Glossary	28

Page numbers refer to content pages (excluding cover and table of contents).

1. What is ClaudeOS?

ClaudeOS is an AI-powered operating system built to help you and your team get more done -- faster and smarter. Think of it as a highly capable AI assistant that never sleeps, remembers everything you tell it, automatically runs tasks on a schedule, and organises all of its work neatly for you to find later.

Unlike a simple chatbot, ClaudeOS has multiple specialised AI agents (think: departments), a memory system (think: a filing cabinet that grows smarter over time), and automated workflows (think: recurring tasks that run themselves). Everything is controlled from a web dashboard you open in your browser.

What Problems Does It Solve?

- You need to write a client report every Friday -- ClaudeOS does it automatically.
- You want to research a topic and get a summary -- one click and it is done.
- You need to remember a client preference or decision -- save it to memory and every future AI response will know about it.
- You want to track all AI-generated content in one place -- every output is saved and searchable.
- You manage multiple clients and do not want their data to mix -- each client has their own private workspace.

The 7 Core Modules (Plain English)

Module	What It Is
Dashboard	A web page you open in your browser to c
API	The behind-the-scenes engine the dashboa
Memory Engine	A knowledge store the AI reads before ev
Agent Registry	12 specialised AI workers, each with a s
Workflow Engine	Automated pipelines that run on a schedu
Client Vault	Private workspaces for each client -- da
Output Manager	Saves and organises every piece of conte

Who Uses ClaudeOS?

ClaudeOS is designed for two types of users:

- End users -- clients and team members who use the dashboard to run agents, search memory, view outputs, and monitor workflows. No technical knowledge needed.
- Administrators -- the person who sets up and maintains the system. They handle configuration, API keys, and server startup. Technical reference is in Section 12 of this guide.

2. Design Principles

ClaudeOS was built with 8 guiding principles that shape how everything works. You do not need to memorise these -- but understanding them will help you get the most out of the system.

1. Windows-First

ClaudeOS runs natively on Windows 10 and above. No Linux server, no cloud subscription required just to run it. Start it with a single PowerShell script and it is ready in under 30 seconds. Everything is designed to work reliably on the Windows desktop.

2. Client Privacy (Namespace Isolation)

Every client has their own private workspace called a namespace. When an AI agent works on a RECI Transport task, it can only see RECI Transport data. Ivycandy Hair data stays in the Ivycandy Hair namespace. There is no risk of one client's information appearing in another client's output.

3. Smart Memory (Dual-Index)

ClaudeOS has two search engines working together. The first is keyword search (like searching a document for an exact word). The second is meaning-based search (finds results even when you use different words). Together they make sure you always find what you are looking for in memory.

4. Background Execution

When you dispatch an AI agent, ClaudeOS runs it in the background immediately. You get a run ID and can continue using the dashboard. No waiting, no frozen screens. Check back any time to see results.

5. Workflow Automation

Workflows are sequences of AI agents that run one after another. Some run automatically on a schedule (e.g. every weekday morning). Others you trigger with one click. You can pass information between steps -- the output of one agent becomes the input of the next.

6. Auto-Tagging

Every time the AI produces a document, report, or draft, ClaudeOS automatically reads the content and assigns relevant tags (like "client", "urgent", "report"). This means you can find any output months later without remembering exactly what it was called.

7. Cloud Backup

ClaudeOS can automatically push your data to Supabase (a secure cloud database) every 15 minutes. This is optional but recommended. It means your memory entries, agent runs, and outputs are backed up safely even if something happens to your local machine.

8. Voice Matching

The Writing and Communications agents adapt their tone based on who the client is. If you have saved context about a client preferring formal emails, the agent will write formally. If a client prefers a warm and casual tone, the agent matches that. This context lives in the client's memory namespace.

3. Getting Started

Prerequisites

Before starting ClaudeOS for the first time, make sure the following are in place:

- Python 3.11 or newer installed on your Windows machine
- An Anthropic API key (this is what powers the AI -- your administrator will have this)
- A .env file in the ClaudeOS folder with your configuration (see Section 12)
- All Python packages installed (your administrator runs this once: `pip install -r requirements.txt`)

Starting the System

ClaudeOS starts with a single command. Open PowerShell, navigate to the ClaudeOS folder, and run:

```
.\scripts\start.ps1
```

The startup script does the following automatically:

1. Reads your port settings from the .env file
2. Stops any old processes that might be running on those ports
3. Runs any pending database updates (migrations)
4. Starts the Flask API server in the background on port 5000
5. Waits and checks that the API is healthy (up to 10 retries)
6. Starts the Streamlit dashboard on port 8501

SUCCESS INDICATOR

When startup is complete you will see "ClaudeOS v1.0.0 - Started" in green. The terminal stays open showing the dashboard. Press Ctrl+C to stop everything.

Opening the Dashboard

Once started, open your web browser and go to:

```
http://localhost:8501
```

You will see the ClaudeOS dashboard with a green sidebar. The sidebar shows CLAUDEOS at the top, the navigation menu, an API status indicator (green dot = online), and the current date and time.

Dark Mode and Light Mode

ClaudeOS supports both dark and light colour themes. At the bottom of the sidebar you will see a toggle button: "Light mode" or "Dark mode". Click it to switch. Your preference is remembered for the session.

API Access (for developers)

The REST API runs separately and can be accessed directly:

```
http://localhost:5000/api/v1/health
```

All API requests require an API key header. Your administrator sets this up in the `.env` file. See Section 12 for full API reference.

4. Dashboard Walkthrough

The ClaudeOS dashboard has 7 pages, each accessible from the sidebar. Click any item in the navigation list to switch pages instantly.

4.1 Overview Page

Purpose: Your home screen -- system health, live activity, and quick dispatch.

When to use: When you open the dashboard, this is the first page you see.

What You See

- Five KPI counters at the top: total memory entries, agents, agent runs, workflows, and outputs
- System Status panel -- shows whether the database, API, and vector search are all running (green = healthy)
- Recent Events -- a live list of the last 8 agent runs with status icons
- Quick Dispatch -- a form on the right where you can run any agent immediately
- Memory by Namespace -- shows how many memory entries each client workspace has

What You Do (Step-by-Step)

- Use Quick Dispatch: select an agent from the dropdown, choose a namespace (client), type your prompt, and click "Run Agent"
- You will receive a Run ID. You can check run results on the Agents page.
- Watch Recent Events to see if your agent run appears with a green tick (done) or red cross (failed)

Expected Output

Agent starts running in the background. A success message shows the Run ID.

4.2 Agents Page

Purpose: Browse, search, and dispatch all 12 AI agents.

When to use: Use this page when you need to run a specific agent and want to see its full details first.

What You See

- Search bar -- type a word to filter agents by name, description, category, or tag
- Category filter -- dropdown to show only agents of one type (e.g. research, content)
- Agent cards -- each card shows the agent name, category badge, description, model, and token limit
- Dispatch Panel -- at the bottom: select agent, namespace, write prompt, toggle "Save output", and dispatch
- Run Status Checker -- after dispatching, click "Check run status" to see results
- Recent Runs Table -- last 20 agent runs with status, tokens used, and duration

What You Do (Step-by-Step)

- Use the search bar to find the right agent for your task
- In the Dispatch Panel, select the agent, choose the correct client namespace, write your prompt, and click "Dispatch"
- If "Save output" is ticked, the AI's response will be saved to the Outputs page automatically

Expected Output

A run starts in the background. The run ID appears. Use "Check run status" to see the response.

4.3 Memory Page

Purpose:	Search, browse, and add knowledge entries that the AI remembers.
When to use:	Use this page to add information about a client, search for a fact you stored, or review all saved knowledge.

What You See

- Stats bar -- shows total memory entries and a count per namespace
- Search tab -- search by keyword, meaning, or both. Select a namespace to narrow down
- Browse tab -- filter by namespace, category, and minimum confidence level
- Add Entry tab -- a form to add a new memory entry manually
- Import button -- imports memory files from the Claude Code memory folder

What You Do (Step-by-Step)

- To search: go to the Search tab, type your query, choose a mode (text / semantic / both), and optionally select a namespace
- To add a memory: go to Add Entry, fill in the namespace, category, key (short label), value (the actual information), optional tags, and confidence level (0.0 to 1.0)
- To browse: use the Browse tab to see all entries for a client

Expected Output

Search returns matching entries with category badge, confidence score, and tags. Add Entry saves the fact immediately.

4.4 Workflows Page

Purpose:	View, trigger, and manage automated multi-step pipelines.
When to use:	Use this page to manually trigger a workflow, enable/disable scheduled workflows, or check past run history.

What You See

- Workflows tab -- lists all 7 workflows with trigger type icon (clock = scheduled, play = manual)

- Each workflow can be expanded to see its description, trigger schedule, and a "Run now" button
- Run History tab -- all recent workflow runs with status, duration, and step details
- Scheduler tab -- shows the next scheduled run time for automatic workflows

What You Do (Step-by-Step)

- To trigger a workflow manually: click on a workflow to expand it, select a namespace, optionally add context (key=value pairs), and click the Run button
- To enable or disable a workflow: click the Enable/Disable button inside the expanded workflow
- To view results: go to Run History, find your run, and click to expand the step log

Expected Output

Workflow starts running. A run ID is returned. Check Run History for step-by-step results and final output.

4.5 Projects Page (Client Vault)

Purpose: Manage client namespaces, projects, and context files.

When to use: Use this page to create or view client workspaces, manage projects within them, and upload context files.

What You See

- Namespaces tab -- lists all client workspaces with icon, name, type, and workspace statistics
- Projects tab -- all projects grouped by client namespace, with status and priority
- Context Files tab -- manage per-namespace context files that agents read for background information
- New Namespace form -- create a new client workspace with name, description, colour, and icon
- New Project form -- add a project under an existing namespace

What You Do (Step-by-Step)

- To view a namespace: click to expand and see workspace file stats and project count
- To add context: go to Context Files, select the namespace, enter a filename and content, and click Save
- To create a project: use the New Project form, fill in name, slug, tech stack, and priority

Expected Output

New namespaces and projects appear immediately. Context files are saved to the vault and agents will read them on next run.

4.6 Outputs Page

Purpose: Browse, search, and download all AI-generated content.

When to use: Use this page to find any report, draft, or analysis the AI has produced.

What You See

- Browse tab -- filter outputs by namespace, type (report/draft/analysis/code/note/archive), and quantity
- Each output shows title, namespace, date, size, tags, and summary
- Actions per output: View full content, download as Markdown file, or delete
- Search tab -- full-text search across all outputs by keyword
- Stats tab -- output counts broken down by type and namespace

What You Do (Step-by-Step)

- To find an output: use the Browse tab and set your filters, or go to Search and type keywords
- To read content: click "View" inside an output to expand the full text
- To download: click the Markdown download link next to the output
- To delete: click the Delete button and confirm

Expected Output

Outputs display with metadata. Download delivers a .md file. Delete removes from database and file system.

4.7 Settings Page

Purpose: Monitor and control Supabase cloud sync.

When to use: Use this page to check whether cloud backup is configured and to manually trigger a sync.

What You See

- Connection badge -- green "Connected" or red "Not Configured"
- Summary metrics -- total rows pushed, failed rows, and auto-sync interval
- Manual Controls -- "Push All Now" button and per-table push option
- Table Sync State -- last sync time and push/fail counts per database table
- Sync Log -- history of the last 30 sync operations
- SQL Schema helper -- Supabase schema SQL for initial setup

What You Do (Step-by-Step)

- To check sync status: just open the page -- the badge tells you immediately
- To trigger a manual sync: click "Push All Now" and wait for the success message
- If not configured: follow the instructions to add SUPABASE_URL and SUPABASE_SERVICE_KEY to your .env file

Expected Output

Push All Now syncs all tables to Supabase. Results show rows pushed and any failures.

5. AI Agents Reference

ClaudeOS has 12 specialised AI agents. Each agent is built for a specific type of task. Below you will find a plain-English description of every agent, what input it needs from you, and what it produces.

HOW TO USE AN AGENT

Go to the Agents page (or the Quick Dispatch panel on the Overview page). Select the agent, choose the correct client namespace, write your prompt, and click Dispatch.

Morning Briefing Agent

Category:	Operations
What it does:	Generates a structured daily briefing by reading recent memory and activity in the namespace.
What to type:	Usually runs automatically. If triggered manually, your prompt can specify focus areas (e.g. "Focus on RECI fleet issues today").
Output format:	5 sections: Priority Actions, Project Status, Client Updates, Scheduled Work, Notes.
Best used for:	Every weekday morning when you want a summary of what needs attention today.

Research Agent

Category:	Research
What it does:	Conducts deep, structured research on any topic. Distinguishes between high-confidence facts and uncertain claims.
What to type:	A clear research question or topic. Example: "Research electric vehicle adoption trends in Nigeria 2024".
Output format:	6 sections: Overview, Key Facts, Key Players, Implications, Sources, Next Actions.
Best used for:	When you need thorough background information on a topic, market, or decision.

Writing Agent

Category:	Content
What it does:	Drafts professional written content: reports, emails, proposals, blog posts. Adapts tone to match the namespace voice.
What to type:	Type of document, purpose, and key points to include. Example: "Write a 1-page project proposal for installing GPS trackers on RECI vehicles."
Output format:	A complete, ready-to-edit document in the requested format.

Best used for: When you need a first draft of any written content -- the agent does 80% of the work, you review and polish.

Communications Agent

Category: Communications

What it does: Drafts client-facing messages: emails, WhatsApp messages, proposals. Matches the client's preferred communication style from memory.

What to type: Who the message is for, what you need to say, and any relevant background. Example: "Draft a follow-up email to the RECI Transport client about the overdue payment."

Output format: A ready-to-send message with subject line (for emails) and appropriate tone.

Best used for: When you need to communicate with a client and want the AI to draft it in the right tone.

Analysis Agent

Category: Analysis

What it does: Analyses data, situations, or decisions. Produces structured findings, patterns, risks, and ranked recommendations.

What to type: What you want analysed and any data or context. Example: "Analyse the booking trends from RECI Transport over the last month and identify the busiest routes."

Output format: 6 sections: Summary, Key Findings, Anomalies, Trends, Recommendations (ranked), Confidence Level.

Best used for: When you need to make sense of data or want a structured breakdown of a complex situation.

Client Manager Agent

Category: Operations

What it does: Generates per-client status cards showing active projects, milestones, blockers, and immediate action items.

What to type: A prompt describing which client or all clients. Example: "Give me a status summary for all active clients."

Output format: One status card per client with: Active Projects, Next Milestone, Current Blockers, Required Actions.

Best used for: For weekly reviews, client meetings, or when you need a quick overview of all client work.

Scheduling Agent

Category:	Operations
What it does:	Converts natural language scheduling requests into structured reminders saved to memory with expiry dates.
What to type:	A plain English scheduling request. Example: "Remind me to follow up with Ivycandy Hair on their wig delivery on Friday at 3pm."
Output format:	A confirmation message plus a memory entry saved with the correct expiry date.
Best used for:	When you want to schedule a reminder or follow-up without leaving the dashboard.

RECI Transport Ops Agent

Category:	Domain (RECI Transport only)
What it does:	Specialised analysis for RECI Transport Ltd. Reviews fleet utilisation, booking trends, revenue, and operational issues. This agent is locked to the reci-transport namespace and cannot run for other clients.
What to type:	A specific operational question or "Run a full fleet analysis." Example: "Which routes had the most cancellations this week?"
Output format:	5 sections: Fleet Utilisation, Booking Trends, Revenue Analysis, Anomalies, Action Items.
Best used for:	For transport-specific analysis, fleet reports, or operational decision-making for RECI Transport.

Meta Orchestrator Agent

Category:	System
What it does:	Breaks down a complex goal into an ordered sequence of agent tasks. Produces an execution plan that can be used to run multiple agents step by step.
What to type:	A high-level goal or complex task. Example: "Plan a full content strategy for Ivycandy Hair's social media launch next month."
Output format:	A JSON execution plan with steps, agents, prompts, and dependencies. Then used by the Meta Orchestration workflow.
Best used for:	When your task is too complex for one agent and needs multiple agents working in sequence.

Memory Curator Agent

Category:	System
What it does:	Reviews all memory entries in a namespace for duplicates, contradictions,

and stale information. Recommends merges and expirations.

What to type:	Usually runs automatically (Sunday midnight). If triggered manually, specify the namespace.
Output format:	A list of memory recommendations: entries to merge, expire, or keep. Plus a summary.
Best used for:	After many memory entries accumulate, or when search results seem noisy or contradictory.

Workflow Builder Agent

Category:	System
What it does:	Converts a plain English description of an automation into a valid ClaudeOS workflow YAML file.
What to type:	A description of the automation you want. Example: "Every Monday morning, research the top 3 AI news stories and email me a summary."
Output format:	A complete workflow YAML definition ready to save and use.
Best used for:	When you want to create a new automated pipeline without writing YAML yourself.

QA Agent

Category:	Engineering
What it does:	Reviews code or outputs for quality, security issues (OWASP), Windows compatibility, and brand compliance.
What to type:	The code or content to review. Example: "Review this Python script for security issues."
Output format:	A markdown table: Severity File/Item Issue Recommended Fix.
Best used for:	Before deploying any code changes, or when reviewing AI-generated code for correctness.

6. Workflows Reference

Workflows are automated pipelines -- sequences of AI agents that run one after another. Some run automatically on a schedule. Others you trigger manually with one click.

HOW TO TRIGGER MANUALLY

Go to the Workflows page. Click on any workflow to expand it. Select a namespace, optionally add context variables (key=value), and click the Run button.

6.1 Morning Briefing

Schedule:	Monday to Friday, automatically at 07:00 WAT (West Africa Time)
What it does:	Reads all high-confidence memory entries in the chosen namespace, reviews recent agent runs, and produces a structured daily briefing covering what needs attention today.
Inputs needed:	None required for automatic runs. For manual: optionally add context like "focus=fleet issues".
Output:	A saved report with 5 sections: Priority Actions, Project Status, Client Updates, Scheduled Work, Notes. Appears in the Outputs page under type "report".

6.2 Memory Curation

Schedule:	Automatically every Sunday at midnight WAT
What it does:	Reviews all memory entries for a namespace, identifies duplicates, contradictions, and stale information, and produces a set of recommendations for cleaning up the knowledge base.
Inputs needed:	None. Runs automatically.
Output:	A saved report listing which memory entries to merge, expire, or keep. The curator does not delete anything automatically -- a human reviews the recommendations.

6.3 Research Digest

Schedule:	Manual -- triggered by you when needed
What it does:	Two-step pipeline: first the Research Agent does a deep dive on your topic; then the Writing Agent condenses the findings into a concise executive summary (around 300 words).
Inputs needed:	topic -- the subject to research. Example: topic=solar energy trends in West Africa context_notes -- optional extra background. Example:

	context_notes=focus on off-grid solutions
Output:	A saved executive summary with key findings, implications, and 3 recommended next actions.

6.4 Client Weekly Report

Schedule:	Automatically every Friday at 17:00 WAT
What it does:	Two-step pipeline: first the Research Agent gathers a week's worth of activity, decisions, and outputs for the namespace; then the Writing Agent formats a professional status report.
Inputs needed:	None. Runs automatically for each client namespace.
Output:	A formatted weekly status report covering completed work, current blockers, and next steps. Saved to Outputs and ready to send to the client.

6.5 Analysis Run

Schedule:	Manual -- triggered by you when needed
What it does:	Two-step pipeline: first the Analysis Agent produces a structured breakdown of the subject; then the Writing Agent reformats it into a clean, readable report under 500 words.
Inputs needed:	subject -- what to analyse. Example: subject=Q3 booking data for RECI Transport data_summary -- any numbers or data to include analysis_goal -- what decision or question this analysis supports
Output:	A structured analysis report with findings, patterns, risks, and ranked recommendations.

6.6 QA Sweep

Schedule:	Automatically Monday to Friday at 06:00 WAT (before the morning briefing)
What it does:	Reviews the previous day's agent runs and outputs for quality issues, errors, anomalies, and anything that did not meet expected standards.
Inputs needed:	None. Runs automatically.
Output:	A pass/fail verdict for yesterday's activity, a list of flagged items, and recommended actions. Saved to Outputs as type "report".

6.7 Meta Orchestration

Schedule:	Manual -- triggered when you have a complex, multi-step task
What it does:	Three-step pipeline: the Meta Orchestrator plans the task; the Research Agent gathers context; the Writing Agent synthesises a final action brief combining the plan and research.
Inputs needed:	task -- the complex goal. Example: task=Build a content launch plan for Ivycandy Hair's new wig range namespace -- which client constraints -- any limits (e.g. constraints=budget under 50k, launch in 3 weeks)
Output:	A comprehensive action brief combining an execution plan, supporting research, and clear next steps.

7. Client Workspace Sections

ClaudeOS uses namespaces to keep each client's data completely separate. Below is a guide for each client workspace -- which agents are most relevant, which workflows apply, and example prompts to get you started.

7.1 RECI Transport (reci-transport)

RECI Transport Ltd is a fleet and logistics client. ClaudeOS manages operational analysis, booking intelligence, fleet reporting, and weekly client status reports for this workspace.

Most Useful Agents

- RECI Transport Ops Agent -- the only agent locked exclusively to this namespace. Use for fleet utilisation, route analysis, and booking trends.
- Analysis Agent -- for interpreting booking data, revenue patterns, or driver performance
- Writing Agent -- for drafting operational reports or client communications
- Communications Agent -- for client-facing emails and follow-ups
- Briefing Agent -- for a daily summary of RECI-specific activity

Scheduled Workflows (Automatic)

- Morning Briefing -- weekdays at 07:00 WAT: daily RECI activity summary
- Client Weekly Report -- Fridays at 17:00 WAT: weekly status report for RECI
- QA Sweep -- weekdays at 06:00 WAT: quality check on yesterday's outputs

Example Prompts

Transport Ops Agent: "Analyse the booking data from this week. Which routes had the highest cancellation rate and why?"

Writing Agent: "Write a 1-page monthly fleet performance summary for the RECI Transport client."

Communications Agent: "Draft a professional follow-up email to the client reminding them about the pending invoice."

Memory Tips for RECI Transport

- Save client contact preferences as "preference" category entries (e.g. key: email_tone, value: formal)
- Save fleet facts as "fact" entries (e.g. key: total_vehicles, value: 12 trucks + 4 vans)
- Save important decisions as "decision" entries (e.g. key: gps_install_decision, value: approved by director 15 May 2026)

7.2 Ivycandy Hair (ivycandy-hair)

Ivycandy Hair is a luxury wig and hair brand. ClaudeOS supports brand communications, content creation, client engagement, and business intelligence for this workspace. Note: Ivycandy Hair also has a separate order management system (ivycandy-hair-process) that handles WhatsApp

orders, payment detection, and DHL shipping. ClaudeOS and the order system work independently -- ClaudeOS handles the AI/content side.

Most Useful Agents

- Writing Agent -- for blog posts, product descriptions, Instagram captions, and brand content
- Communications Agent -- for customer emails, WhatsApp templates, and supplier messages
- Research Agent -- for market research, competitor analysis, and beauty industry trends
- Analysis Agent -- for sales performance, customer feedback analysis, and seasonal trends
- Client Manager Agent -- for weekly business overview and project tracking

Scheduled Workflows (Automatic)

- Morning Briefing -- weekdays at 07:00 WAT: daily Ivycandy business summary
- Client Weekly Report -- Fridays: weekly brand performance review
- Memory Curation -- Sundays: keeps knowledge base clean and up to date

Example Prompts

Writing Agent: "Write 5 Instagram caption options for a new Brazilian wavy wig launch. Tone: luxurious, aspirational, Nigerian audience."

Research Agent: "Research the top wig trends for 2026 among Nigerian women aged 25-40."

Communications Agent: "Draft a professional apology email to a customer whose order was delayed by 3 days."

Memory Tips for Ivycandy Hair

- Save brand voice as a "preference" entry (e.g. key: brand_voice, value: warm, aspirational, confident, Nigerian)
- Save product lines as "fact" entries (e.g. key: product_range, value: Brazilian wavy, Bone straight, Kinky curly)
- Save pricing as "context" entries for use in content (e.g. key: price_range, value: N25,000 - N150,000)

7.3 Faiyke AI (faiyke-ai)

Faiyke AI is an AI SaaS client. ClaudeOS supports research, technical writing, product strategy, and business analysis for this workspace.

Most Useful Agents

- Research Agent -- for AI industry research, competitor analysis, and market sizing
- Writing Agent -- for technical documentation, product copy, and pitch materials
- Analysis Agent -- for product metrics, user feedback analysis, and strategic decisions
- Meta Orchestrator -- for complex multi-step product or strategy tasks
- QA Agent -- for reviewing AI-generated code or technical outputs

Example Prompts

Research Agent: "Research the current landscape of AI agent platforms in 2026. Focus on pricing models and target markets."

Writing Agent: "Write a 2-page executive summary for a pitch deck for Faiyke AI, positioning it as an enterprise AI automation platform."

Meta Orchestrator: "Plan a 3-month product launch campaign for Faiyke AI including research, content, and outreach strategy."

7.4 Personal & General (personal / global)

The "personal" and "global" namespaces are for internal use and general tasks not tied to a specific client. Use "global" for system-wide shared information. Use "personal" for your own scheduling, reminders, and research.

Common Uses

- Scheduling -- use the Scheduling Agent to create reminders (e.g. "Remind me to review Q2 accounts on June 1st")
- Personal research -- use the Research Agent for topics you are personally exploring
- System memory -- store facts that apply across all clients (e.g. your own contact details, standard rates)
- Draft templates -- save reusable document templates as memory entries

Scheduling Agent: "Remind me to follow up on the RECI payment on Friday 23rd May at 9am."

Research Agent (personal): "Summarise the key AI tools a digital agency should be using in 2026."

8. Memory System Guide

ClaudeOS memory is like a smart filing cabinet. Every time you save a fact, decision, or preference, it gets stored with a label (key), the content (value), a category, and a confidence score. Before every agent run, ClaudeOS automatically reads the most relevant, high-confidence memory entries and includes them in the AI's context.

This means the more you save to memory, the smarter and more personalised every agent response becomes -- without you having to repeat yourself.

Memory Categories

Category	What It Stores	Example
fact	Objective, verifiable information	key: vehicle_count value: 12 trucks
decision	A choice or approval that was made	key: gps_install value: approved May 20
context	Background information for a topic	key: q3_goals value: expand to Abuja ma
preference	Client or personal preference	key: email_tone value: formal and conci
reminder	A time-based note to act on	key: invoice_followup value: chase RECI
insight	A conclusion or observation	key: peak_hours value: bookings spike F

How to Add a Memory Entry (Step-by-Step)

1. Go to the Memory page from the sidebar
2. Click the "Add Entry" tab
3. Select the Namespace (e.g. reci-transport)
4. Select the Category (e.g. fact)
5. Enter a Key -- a short label for this entry (e.g. fleet_size)
6. Enter a Value -- the actual information (e.g. 12 trucks and 4 vans, as of May 2026)
7. Optionally add Tags (comma-separated, e.g. fleet, vehicles, reci)
8. Set Confidence -- how sure you are this is correct (1.0 = certain, 0.5 = estimated)
9. Click Save

How to Search Memory

ClaudeOS has two search engines. You can use either or both:

- Text search -- finds entries that contain your exact words or similar keywords. Good for: "find that entry I wrote about GPS installers."
- Semantic search -- finds entries by meaning, even if you use completely different words. Good for: you type "vehicle tracking" and it finds an entry about "GPS monitoring."
- Both -- uses text and semantic together, then merges results. Best for thorough searches.

To search:

1. Go to the Memory page, Search tab

2. Type your query
3. Select mode: text, semantic, or both
4. Optionally select a namespace to narrow results
5. Results appear with category badge, confidence score, and matching tags

Confidence Scores Explained

Score	Meaning
1.0	Certain fact -- confirmed and verified
0.9	Very high confidence -- strong source
0.7	Reasonable confidence -- likely correct
0.5	Estimated -- use with caution
Below 0.5	Uncertain -- may need verification

NOTE

Agents only read memory entries with confidence 0.8 or above by default. Low-confidence entries are still stored and searchable, but they do not automatically influence agent responses.

Namespace Isolation

Memory is completely separate per namespace. When you run an agent in the reci-transport namespace, it only reads reci-transport memory. It cannot see ivycandy-hair or personal entries. This ensures client privacy is maintained at all times.

9. Outputs Guide

Every time an AI agent or workflow produces a document, report, draft, or analysis, ClaudeOS saves it automatically. These saved items are called outputs. You can browse, search, download, or delete them from the Outputs page at any time.

Output Types

Type	What It Contains
report	Completed analysis or status reports
draft	Unfinished documents, proposals, or temp
analysis	Data analysis, breakdowns, and findings
code	Any code snippets or scripts produced by
note	Short notes or summaries
archive	Older outputs moved to long-term storage

How to Browse Outputs (Step-by-Step)

1. Go to the Outputs page from the sidebar
2. In the Browse tab, select a Namespace (or leave as "all")
3. Select a Type filter (or leave as "all")
4. Set the Limit (20, 50, or 100 results)
5. Results appear as expandable cards with title, date, size, and tags
6. Click an output card to expand it and see the summary

How to View Full Content

1. Expand the output card
2. Click the "View" button
3. The full content appears below the card
4. Click "Hide" to collapse it again

How to Download an Output

1. Expand the output card
2. Click the "Download MD" link
3. A Markdown (.md) file downloads to your browser's default download folder
4. Open with any text editor, Word (via conversion), or Notion

How to Search Outputs

1. Go to the Outputs page, Search tab
2. Type keywords (e.g. "RECI fleet report June")
3. Optionally filter by namespace

4. Click Search
5. Results show title, namespace, type, and a relevance score

Auto-Tagging

ClaudeOS reads every output and automatically assigns relevant tags. For example, a document mentioning "fleet", "trucks", and "RECI" will be tagged with those words. This makes future searches much easier -- you do not have to remember exact titles.

Deleting an Output

WARNING

Deleted outputs are permanently removed from the database and the file system. There is no undo. Download the output first if you might need it later.

1. Expand the output card
2. Click the "Delete" button
3. A confirmation prompt appears
4. Click "Yes, delete" to confirm or "Cancel" to go back

10. Technical Reference (Administrators)

This section is for the person who installs, configures, and maintains ClaudeOS. End users do not need to read this section.

Environment Configuration (.env File)

ClaudeOS reads all configuration from a .env file in the project root. Create this file before starting the system for the first time.

```
# Core
ANTHROPIC_API_KEY=sk-ant-your-key-here
CLAUDEOS_SECRET_KEY=change-this-to-a-random-string
CLAUDEOS_ENV=development
CLAUDEOS_VERSION=1.0.0

# Ports
FLASK_PORT=5000
STREAMLIT_PORT=8501

# Database
SQLITE_PATH=data/claudeos.db
CHROMADB_PATH=data/chromadb

# Logging
LOG_LEVEL=INFO
LOG_PATH=logs

# Optional: Supabase Cloud Sync
SUPABASE_URL=https://your-project.supabase.co
SUPABASE_SERVICE_KEY=your-service-role-key

# Optional: Dev API key for dashboard auth
CLAUDEOS_DEV_API_KEY=your-dev-api-key
```

Key File Locations

File / Folder	Description
data/claudeos.db	Main SQLite database (all tables)
data/chromadb/	ChromaDB vector embeddings
logs/api.log	API server log (max 5MB, rotates)
vault/workspaces/	Per-namespace context and file storage
outputs/store/	Saved output files organised by type/nam
agents/definitions/	12 YAML agent definition files
workflows/definitions/	7 YAML workflow definition files
.env	Environment configuration (never commit

First-Time Setup Commands

Run these once after cloning the project:

```
pip install -r requirements.txt
python scripts/migrate.py
python scripts/seed_agents.py
python scripts/seed_workflows.py
python scripts/seed_namespaces.py
```

API Quick Reference

Method	Endpoint	Description
GET	/api/v1/health	Check API is running
GET	/api/v1/system/status	Full system health check
GET	/api/v1/agents	List all agents
POST	/api/v1/agents/{name}/run	Run an agent (returns run_id)
GET	/api/v1/agents/runs/{id}	Check agent run status
POST	/api/v1/memory/search	Search memory entries
POST	/api/v1/memory	Create memory entry
GET	/api/v1/workflows	List all workflows
POST	/api/v1/workflows/{name}/run	Trigger workflow
GET	/api/v1/outputs	List outputs
GET	/api/v1/sync/status	Check Supabase sync status
POST	/api/v1/sync/push	Push data to Supabase

Supabase Cloud Sync Setup

1. Create a free Supabase project at supabase.com
2. In Supabase, go to SQL Editor and run the schema SQL from: `sync/supabase_schema.sql`
3. Copy your Project URL and Service Role key from Supabase Settings > API
4. Add `SUPABASE_URL` and `SUPABASE_SERVICE_KEY` to your `.env` file
5. Restart ClaudeOS (`start.ps1`)
6. Open the Settings page in the dashboard -- the badge should show green "Connected"
7. Click "Push All Now" for the first full sync

11. Troubleshooting

Problem: API shows red dot (offline) in sidebar

Possible causes:

- The Flask API is not running
- A previous process is still holding port 5000
- An error occurred during startup

How to fix:

- Open PowerShell and run: `.\scripts\start.ps1`
- Check the terminal for red error messages during startup
- Check logs/api.log for detailed error information
- If port 5000 is blocked: run `netstat -ano | findstr :5000`, then `taskkill /F /PID <number>`

Problem: Dashboard will not open (<http://localhost:8501>)

Possible causes:

- Streamlit has not started yet
- Port 8501 is in use by another process

How to fix:

- Wait 10 seconds after running `start.ps1` and try again
- Check the PowerShell terminal -- Streamlit startup messages should appear
- Kill port 8501: `netstat -ano | findstr :8501` then `taskkill /F /PID <number>`
- Restart with: `.\scripts\start.ps1`

Problem: Agent run shows "failed" status

Possible causes:

- Prompt was empty or too short
- Anthropic API key is missing or expired
- The agent hit its token limit

How to fix:

- Check the error message in the Run Status panel on the Agents page
- Verify `ANTHROPIC_API_KEY` is set in your `.env` file and is valid
- Try a shorter, more focused prompt
- Check logs/api.log for the detailed Python error

Problem: Memory search returns no results

Possible causes:

- No memory entries exist for this namespace yet
- Search terms do not match stored entries
- ChromaDB vector store is not initialised

How to fix:

- Go to the Memory page, Browse tab -- check if any entries exist
- Try switching between text, semantic, and both search modes
- Try broader or different search terms
- Restart the server to reinitialise ChromaDB: `.\scripts\start.ps1`

Problem: Supabase sync shows "Not Configured"**Possible causes:**

- SUPABASE_URL or SUPABASE_SERVICE_KEY is missing from `.env`
- Server has not been restarted after adding credentials

How to fix:

- Open `.env` and verify both SUPABASE_URL and SUPABASE_SERVICE_KEY are present
- Restart the server: `.\scripts\start.ps1`
- Open Settings page and check the connection badge
- If still failing: check your Supabase project is active and the SQL schema has been run

Problem: Workflow did not run automatically**Possible causes:**

- Workflow is disabled
- APScheduler did not start correctly
- Server was not running at the scheduled time

How to fix:

- Go to Workflows page and check if the workflow shows as enabled (green dot)
- Check the Scheduler tab -- does the workflow appear with a next run time?
- Click "Reload Scheduler from DB" in the Scheduler tab
- Ensure the server runs continuously if you rely on scheduled workflows

12. Glossary

Key terms used in this guide, explained in plain English.

Agent	A specialised AI worker with a specific job. ClaudeOS has 12 agents, each designed for a different task type.
API	Application Programming Interface. The behind-the-scenes engine that the dashboard talks to. Runs on port 5000.
APScheduler	The scheduling engine that runs workflows automatically at the right time. Starts with the server.
ChromaDB	The vector database that powers semantic (meaning-based) memory search.
Confidence Score	A number from 0.0 to 1.0 indicating how reliable a memory entry is. Agents only read entries above 0.8 by default.
Dispatch	The action of sending a task to an AI agent. You dispatch an agent by writing a prompt and clicking Run.
FTS5	Full-Text Search version 5. The SQLite keyword search engine used for memory and output text search.
Namespace	A private workspace for a specific client or project. Data does not cross between namespaces.
Memory Entry	A saved piece of knowledge with a key, value, category, and confidence score. Agents read these automatically.
Output	Any document, report, draft, or analysis produced by an AI agent and saved by ClaudeOS.
PowerShell	The Windows command-line tool used to start ClaudeOS. Run it by right-clicking and selecting "Run as administrator".
Run ID	A unique identifier assigned to each agent or workflow run. Use it to check the status and results.
Semantic Search	Searching by meaning rather than exact words. Finds relevant results even when you phrase things differently.
SQLite	The local database file (claudeos.db) where all data is stored on disk.
Streamlit	The Python framework that powers the ClaudeOS web dashboard. Runs on port 8501.
Supabase	A cloud database service used for optional backup and sync. Configured in Settings page.
WSGI	Web Server Gateway Interface. ClaudeOS uses "waitress" as its WSGI server (Windows-compatible, no Gunicorn needed).
Waitress	The Windows-compatible web server that hosts the Flask API. Starts automatically via start.ps1.
Watermark	The last-synced timestamp used by Supabase sync to only push new data, not

re-send everything.

Workflow

An automated pipeline of one or more AI agents that run in sequence. Can be scheduled or manually triggered.